Empirical Evaluation of Management Practices I
Predictions and Machine Learning

Prof. Dr. Dirk Sliwka   Jesper Armouti-Hansen

17/11/20

Content

1. Introduction to Machine Learning

General definition

[Machine Learning is the] field of study that gives computers the ability to learn without being explicitly programmed.
- (Arthur Samuel, 1959)

More specific definition

A computer program is said to learn from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.
- (Tom Mitchell, 1997)

Examples of ML/AI in business

- Employee attrition prediction
- Recruiting automation
- Customer churn modeling
- Recommendation engines

Types of Machine Learning

Supervised Learning:

The task of learning a function that maps the input(s) X to an output y based on example input-output pairs.

- Regression: The output is continuous or discrete and ordered.
  - Example: Predicting house prices based on house characteristics.
- Classification: The output is a discrete and unordered set.
  - Example: Classifying an email as spam or ham based on the use of certain words.

This is a "mini-course" on supervised learning with Python.

Unsupervised Learning:

We observe inputs but no output. We can seek to understand the relationship between the variables or the observations.

- For example, we might observe multiple characteristics for potential customers.
  - We can then try to cluster potential customers into groups based on these characteristics
- We might also try to project our inputs into a lower dimensional space.
  - This can be a beneficial preprocessing for supervised learning when dealing with high-dimensional data.

We will not deal with unsupervised learning in this course.

Terminology

*The terms used in the ML literature differs slightly from that used in Econometrics:*

- Supervised learning $\rightarrow$ *Regression, classification, predicting y_i given x_i.*
- Features $\rightarrow$ x_i, *independent variables, explanatory variables, regressors, predictors.*
- Target $\rightarrow$ y_i, *dependent variable.*
- Training $\rightarrow$ *Estimating a model.*
- Testing $\rightarrow$ *Evaluating a model.*
- Training data $\rightarrow$ *The sample we use to train our model.*
- Test data $\rightarrow$ *The sample we use to test our model.*

2. Regression

Let us start with an example:

- Suppose our task is to predict income based on a person's age and that we had data on the whole population:

Population of income based on age

- What would be a "good" prediction here?

More formally

- Let X_i be a random vector (i.e. the vector of features).
- Let Y_i be a real variable (i.e. the response).
- We are interested in a function $f(X_i)$ which makes "good" prediction about Y_i.
- To know what a "good" prediction is, we require a loss function: $L(Y_i, f(X_i))$ which penalizes bad predictions
  - Common choice: Squared error – penalizes the quadratic distance:
$$L(Y_i, f(X_i)) = (Y_i - f(X_i))^2 \tag{1}$$

Recall from the lecture on Part 2:

CEF Prediction Property

Let $f(X_i)$ be any function of X_i. The conditional expectation function (CEF) solves:

$$E[Y_i|X_i] = \arg \min_{f(X_i)} E[(Y_i - f(X_i))^2] \qquad (2)$$

- Thus, in terms of prediction, we can do no better than the CEF
- Also, recall:

### CEF Decomposition Property

We can decompose Y_i such that

$$Y_i = E[Y_i|X_i] + \epsilon_i \qquad (3)$$

Where:

1. ϵ_i is mean independent of X_i: $E[\epsilon_i|X_i] = 0$.
2. ϵ_i is uncorrelated with any function of X_i.

Thus, if we knew the population, calculating the CEF is usually not a difficult task:

Population of income based on age

We just predict the average income for people of the same age.

- In this case, we are dealing with simulated data with the following conditional expectation function:

$$f(income) = 2000 * age - 20 * age^2 \qquad (4)$$

Usually, we do not know the whole population. Rather, we are working with a sample.

- In this example, we will be working with a sample of 100 data points.
- Our goal is to construct a model $\hat{f}(age)$ which makes "good" predictions about income

Estimating f - K Nearest Neighbor (KNN) Regression

The CEF (f) is almost always unknown, so how can we estimate it?

- For a given observation x_i, we could approximate the CEF by predicting the average of Y_i across observations with $X_i = x_i$.
- Problem?
  - We might have very few or no other observations with $X_i = x_i$
- Instead, we can settle with predicting the average of Y_i of the K nearest known neighbors of x_i:

$$\hat{f}(x_i) = \frac{1}{K} \sum_{j \in N_K} y_j \tag{5}$$

  - Where N_K is a neighborhood containing the indices of the K closest x's.

Estimating f - Linear Regression

Another way to estimate the CEF is to assume that it is approximately linear in its arguments:

$$\hat{f}(x_i) = \hat{\beta}_0 + \hat{\beta}_1 x_{i1} + \hat{\beta}_2 x_{i2} + \cdots + \hat{\beta}_k x_{ik} \tag{6}$$

- Thus, estimating the CEF amounts to finding the β's that minimize the squared error in the sample.

- *We now know of two machine learning models: KNN and Linear Regression!*

Example

Let us return to our example:

Underfitting vs. Overfitting

In order to make predictions based on our sample, suppose we train two models: A simple linear regression:

1. $\hat{f}(age_i) = \hat{\beta}_0 + \hat{\beta}_1 age_i$
2. A KNN regression with $K = 1$

Clearly, the linear regression is too inflexible and the 1NN regression is too flexible to properly approximate the CEF.

The simple linear regression makes bad predictions because it underfits our training data:

- It has high bias and low variance.

The 1NN regression makes bad predictions because it overfits our training data:

- It has low bias and high variance.

Thus, when dealing with predictions, we face a trade-off:

- The Bias-Variance trade-off

Essentially, we want to find a model that has low bias and low variance, if possible.

The Bias-Variance trade-off

More formally, given a realized point x_0, the expected squared error is given by:

$$E[(y_0 - \hat{f}(x_0))^2] = Var[\hat{f}(x_0)] + [f(x_0) - E[\hat{f}(x_0)]]^2 + \sigma_\epsilon^2 \quad (7)$$

The first term is the variance of our model at x_0, the second the squared bias at x_0, and the third is the irreducible error.

Introduction to Sci-kit Learn (sklearn)

Linear Regression in sklearn

- To fit a linear regression model on the data X and y, we first import the module:
  ```
  from sklearn.linear_model import LinearRegression
  ```
- Then we can perform a regression with the following code:
  ```
  reg = LinearRegression().fit(X,y)
  ```
- We can access one of its attributes to get the coefficients:
  ```
  reg.coef_
  ```
- Suppose the regression is given by
  $\hat{f}(x_i) = \hat{\beta}_0 + \hat{\beta}_1 * age + \hat{\beta}_2 * age^2$ and that I would like to get the prediction for a 20 year old person:
  ```
  reg.predict([[20, 20**2]])[0]
  ```

KNN Regression in sklearn

- To fit a KNN regression model on the data X and y, we first import the module:
  ```
  from sklearn.neighbors import KNeighborsRegressor
  ```
- Then we can perform a regression with the following code:
  ```
  knn = KneighborsRegressor().fit(X,y)
  ```
- Again, suppose we want the prediction of a 20 year old person:
  ```
  knn.predict([[20]])[0]
  ```

Drawing a random sample with pandas

- Suppose we would like a random sample of 50 observations from a pandas dataframe df:
  ```
  df_sample = df.sample(n=50, random_state=181)
  ```

Your tasks

1. Import the income dataset as a pandas dataframe. It is located at https://raw.githubusercontent.com/armoutihansen/ EEMP2020/main/datasets/income.csv
2. Create a new column in the dataframe "age_sq" that takes the squared values of the age variable.
3. Since the CEF can now be estimated by a linear regression, draw a random sample of 100 observations and fit a linear regression on the sample. Are the coefficients close to those given by the CEF?
4. Write a loop that generates a new sample 100 times of 100 observations and fit a linear regression on each of the 100 samples. Store the coefficients of the model in each loop.
5. Calculate the mean of each coefficient. Are the means close to the coefficients given in the CEF?

6. Consider the observation *age* = 50. What is the conditional expectation of income at that age?
7. Drop the "age_sq" variable from the dataframe and write a loop that generates a new sample 100 times of 100 observations. On each sample, fit a 1NN regression and store its income prediction of a person aged 50 in each sample.
8. Calculate the mean of the predictions and plot the distribution of the predictions. What does this tell you about the 1NN's bias and variance on this dataset?